AWS

SIMPLE STORAGE SERVICE

OR S3 FOR NORMAL PEOPLE

BYGK

WHAT IS S3?

- Object storage
- Designed to be storage for the internet.
- Data can be uploaded and downloaded from S3.
- Almost unlimited amount of data, accessible from anywhere
- 99.999999999% durability (that's eleven nines!)
- Cheapest way to store data on AWS
- Can even host static websites
- Supports BitTorrent, too
- Integrates with many AWS services
- 5TB can be stored in S3

COMMON USE CASES

- Backup and recovery
- Data archiving
- Data lakes
- Hybrid cloud storage
- Cloud-native application data

S3 INTO THREE NAMESPACE

- Namespace is divided into entities called Bucket and Keys.
- A bucket is conceptual container for objects stored in S3
- Bucket Namespace is shared among all the users. This means when creating a bucket , if another user trying to create same bucket it will be not created.

Eg:

S3.amazonaws.com/bucketName/objectName

Global Part

Custom Part

S3 PRICING

- 1) Storage Pricing: Monthly charged for the data stored in S3
- 2) Request Pricing: Amazon S3 charges for each request to the buckets and objects that you own
- 3) Data Transfer Pricing: S3 charges for all the data that is served out of Amazon S3.

BUCKET – CONTAINER RESOURCE

- Logical resource, similar to directory
- Region-specific, but has globally unique name
- Has its own set of access policies and ACLs
- Has multiple bucket-wide options:
  - Versioning
  - Logging
  - Notifications
  - Cross-region replication
  - And many more

OBJECT – KEY-VALUE RESOURCE

- Object is a key-value pair: key is file name, value is the content
- Can be versioned
- Metadata is a set of key-value pairs that store information about an object
- Has subresources, such as torrent and ACL
- Each object has a storage class associated with it

STORAGE CLASSES

- Standard (STANDARD & RRS) – default storage class
  - STANDARD – millisecond access times, full durability/availability
  - RRS – reduced redundancy storage – is meaningless now, don't use it
- Infrequent access (STANDARD_IA & ONEZONE_IA) – for infrequently accessed files
  - STANDARD_IA: millisecond access times, cheaper storage, expensive requests
  - ONEZONE_IA: like standard, but less available/resilient, so its somewhat cheaper
  - Suitable for files over 128KB that you plan to store for at least 30 days
- Glacier – for archiving data
  - Not available in real time! You need to restore objects first
  - Very cheap storage, very expensive requests

STORAGE CLASSES – IN NUMBERS

Storage class	Durability	Availability
STANDARD	99.999999999%	99.99%
RRS	99.99%	99.99%
STANDARD_IA	99.999999999%	99.9%
ONEZONE_IA	99.999999999%	99.5%
GLACIER	99.999999999%	99.99%

VERSIONING

- Off by default
- Useful to prevent unintended deletions or overwrites
- Once versioning is enabled, you cannot disable it (you can still suspend it)
- Each object version is stored separately (takes more space)
- GET request returns the latest version by default – you can specify version id to get specific version
- DELETE request does not delete all versions, it just puts a delete marker as a current version. You can still permanently delete specific versions of an object

ACL – ACCESS CONTROL LISTS

- A resource-based access policy
- Applies both to buckets and objects, each has an ACL attached as a subresource
- Works on account / group level
- Can be used to grant read/write permissions to other accounts
- Limitations:
  - Cannot be used to grant permissions to IAM users
  - No conditional permissions
  - No deny rules

ACL - GRANTEE

- A Grantee is an entity that receives permissions
- A Grantee could be:
  - An AWS account (identified by a Canonical User Id)
  - A predefined group (represented by a URL):
    - Authenticated Users (http://acs.amazonaws.com/groups/global/AuthenticatedUsers)
    - All Users (http://acs.amazonaws.com/groups/global/AllUsers)
    - Log Delivery (http://acs.amazonaws.com/groups/s3/LogDelivery)

ACL - PERMISSION

- Permissions describe which actions a Grantee is allowed to perform on a resource
- You can grant following permissions:

Permission	When granted on a bucket	When granted on an object
READ	Allows to list objects in the bucket	Allows to read object data and metadata
WRITE	Allows to create, overwrite, delete any object in the bucket	Not applicable
READ_ACP	Allows to read bucket ACL	Allows to read object ACL
WRITE_ACP	Allows to write ACL for the bucket	Allows to write ACL for the object
FULL_CONTROL	Same as all of the above	Same as all of the above

POLICIES – POLICY LANGUAGE

- JSON-based documents
- User policies (IAM) and Bucket policies (S3)
- Policies consist of following sections:
  - Resources: buckets and objects in S3, identified by ARN
  - Actions: for each resource you can define a set of operations that will be allowed or denied
  - Effect: allow or deny
  - Principal: account, user, service, or other entity affected by the policy
  - Condition (optional): lets you specify conditions for when your policy is in effect

POLICIES – USER POLICIES

- You can use IAM user policies to control access to S3 resources
- ACLs, bucket policies, and user policies are all affect S3 resources
- Will be covered in IAM section